

Webservers

Server OS

Contents


- HTTP(S)
- Proxy
- Webserver overview
- Apache
- NginX

HTTP(S)

- When talking about the internet, most people refer to HTTP(S).
 - The internet is obviously a lot more.
 - It does show the importance of the HTTP(S) protocol.
- Hypertext Transfer Protocol (Secure)
- Developed by Tim Berners-Lee
- World Wide Web

HTTP(S)

- Client – server / request – response model
 - Clients: webrowsers
 - GET the content from the webserver
 - e.g. Chrome, Firefox, Edge, Internet Explorer, Safari,...
 - Servers: webrowsers
 - Provide the content to the client
 - e.g. Apache, NginX, IIS,...



HTTP(S)

- Application layer protocol
- Uses TCP in the transport layer: sessions
- HTTP: unsecured
 - Default port 80
- HTTPS: secured
 - Default port 443
 - Uses SSL (older version) or TLS (newer version) as encryption protocol
 - HTTP over SSL or HTTP over TLS
 - Secure: authentication, encryption of communication, protection against man-in-the-middle attacks

HTTP(S)

- Frequent HTTP(S) error messages
 - 4xx client errors:
 - 400: Bad request
 - Request is using invalid syntax
 - e.g. corrupt cookie, faulty browser
 - 401: Unauthorized
 - Request requires authentication which has not been provided or was invalid
 - 403: Forbidden
 - Request is understood by the server but the client does not have (file)permission or it is a prohibited action
 - e.g. on linux server, make sure the user account of the webserver (often www-data) has Read permission on the webfiles
 - e.g. missing index.html file
 - 404: Not Found
 - Requested resource is not available on server

<https://www.digitalocean.com/community/tutorials/how-to-troubleshoot-common-http-error-codes>

HTTP(S)


- Frequent HTTP(S) error messages
 - 5xx server errors:
 - 500: Internal Server error
 - Badly configured server
 - 502: Bad Gateway
 - Request sent to a proxy server, but the proxy server is not receiving a valid response from the backend webserver(s)
 - 503: Service Unavailable
 - Server is overloaded or under maintenance
 - e.g. DDOS attack
 - 504: Gateway Timeout
 - Request sent to a proxy server, but the proxy server is not receiving a valid response from the backend webserver(s) within the allowed time period

<https://www.digitalocean.com/community/tutorials/how-to-troubleshoot-common-http-error-codes>

HTTP(S)

- URL

- Uniform Resource Locator (URI with location of resource/file)
- Scheme : http, https, ftp, mailto, file, data, and irc
- Userinfo: (optional): username
- Host: DNS-name or IP
- Path: Path to resource on server
- Query and fragment: optional



HTTP(S)


- URL
 - HTTPS example

Diagram illustrating the components of a URL (https://john.doe@www.example.com:123/forum/questions/?tag=networking&order=newest#top):

- userinfo**: john.doe@
- host**: www.example.com
- port**: 123
- path**: /forum/questions/
- query**: ?tag=networking&order=newest
- fragment**: #top
- scheme**: https
- authority**: john.doe@www.example.com:123

Proxy

- Forward proxy



Proxy


- Forward proxy
 - Are used for filtering and caching
 - Popular forward proxy software:
 - Apache, squid, IIS (Windows), ...
 - An HTTP request from inside the network to an outside webserver is sent to the proxy server instead of directly to the outside webserver.
 - Direct HTTP requests to the internet are blocked by the firewall, making the proxy server the only gateway to the www.

Proxy

- Forward proxy
 - The proxy can perform filtering:
 - Certain URI's are not forwarded (filtered)
 - The proxy server will get the content from the webserver for the webclient and cache the content.
 - If a new request is sent to the proxy server, it will serve the content directly to the client from its local cache.
 - Speed gain (traffic over lan versus wan)
 - Less bandwidth usage over the wan internet connection

Proxy

- Reverse proxy



Proxy


- Reverse proxy
 - On the other side of the HTTP-communication, webserver can also use proxy servers.
 - Reverse proxy
 - Popular reverse proxy software:
 - Haproxy, Squid, NginX, Varnish...
 - The reverse proxy will also be positioned between the internet and the inside of the LAN, this time the webserver.
 - It will take all incoming GET requests from outside webclients and forward them to the actual webserver and send the response to webclient.

Proxy

- Reverse proxy
 - In doing so, several tasks can be handled by the proxy, diminishing the load on the webservers:
 - Static content might be cached, and delivered directly to the client:
 - Web acceleration
 - Content compression
 - TLS encryption/decryption for HTTPS
 - Authentication
 - Load-balancing
 - Also helps protecting the webservers:
 - Hidden behind the proxy
 - DOS attacks

Webservers overview

- Many different webserver software packages are available, with changing popularity over the years.



<https://news.netcraft.com/archives/category/web-server-survey/>

Webservers overview

- Apache
 - Since 1999
 - Open-source
 - Free
 - Cross-platform: Linux, Unix, BSD, but also Windows
 - Vast majority on Linux
 - Often in LAMP setup: **L**inux + **A**pache + **M**ySQL/**M**ariaDB + **P**HP/**P**ython/**P**earl
 - Initial design aimed towards webserver

Webservers overview

- Apache
 - Characteristics
 - Extensible through extensive library of modules , e.g. PHP
 - Inserted in the work processes
 - Dynamically loaded while running the server
 - Can handle dynamic content internally itself through modules
 - Simpler configuration
 - Directory based configuration through `.htaccess` files
 - Allows decentralized configuration of websites
 - Are read with every request: allows changing configuration of a site without restarting the service
 - Web requests are mostly interpreted as physical requests to files
 - The filesystem hierarchy is the same as the available document tree
 - See also the use of the `.htaccess` files

Apache

- Architecture

- Apache server: httpd

- Answers request and provides the content
 - In itself can only deliver static content

- Modules

- Apache can be extended by adding modules for added functionality
 - e.g. authentication, security, php, compression,...
 - Dynamically loaded
 - Only add modules you need
 - Less resources
 - Smaller attack surface

Apache

- Architecture
 - Virtual Hosts
 - One Apache server can deliver multiple different websites by creating them as virtual hosts
 - Security
 - Several components are provided for securing the website, like access control, authentication and encryption
 - Logging

Apache

- Configuration files
 - Apache uses different plain-text files for its configuration:
 - Apache wide main configuration file:
 - `/etc/httpd/conf/httpd.conf`
 - Additional configuration files:
 - Extra configuration files can be added and will be read by Apache if placed in:
 - `/etc/httpd/conf.d/*.conf`
 - Module configuration files:
 - `/etc/httpd/conf.modules.d/*.conf`

Apache

- Configuration files
 - Example configuration files:
 - Apache provides several example configuration files when installing httpd:
 - `/usr/share/doc/httpd/`
 - `httpd-default.conf`
 - `httpd-vhosts.conf`
 - `proxy-html.conf`

Apache

- Directives

- Rules, instructions and settings for how the webserver should behave.
- General directives: for the whole server
- Scoped directives: for a specific part of the server
 - Directories, files, modules and virtual hosts
 - Default in `httpd.conf` is block all, meaning explicit overrides need to be defined

```
<Directory />  
    AllowOverride none  
    Require all denied  
</Directory>
```

- Follows syntax as in the example above

Apache

- Commonly used directives

- `ServerRoot`
 - top-level directory of your Apache configuration
 - `/etc/httpd`
 - configuration files, log files and also your modules
- `Listen`
 - Default is (port) 80
 - Default all IP addresses of the server
 - Can be a different port or specific IP
 - <https://httpd.apache.org/docs/2.4/bind.html>

Apache

- Commonly used directives
 - Include
 - Which config files to include
 - e.g. `/etc/httpd/conf.d`
 - User / Group
 - Which user and which group to run Apache as
 - Good practice to create a specific user for Apache
 - ServerAdmin
 - Email address where errors are sent to

Apache

- Commonly used directives

- `DocumentRoot`
 - Default directory for the actual web documents
 - By default `/var/www/html`
 - <https://httpd.apache.org/docs/2.4/urlmapping.html>
- `Options`
 - Controls which server features are available for a particular directory
- `AllowOverride`
 - Additional directives can be created per directory
 - `.htaccess` file in the directory itself
 - Often used for scoped directives

Apache

- Apache commands
 - `apachectl`
 - The main command used for managing Apache
 - `stop`
 - `start`
 - `restart`
 - `status`
 - `configtest`
 - validate configuration files
 - Often just a wrapper for the `systemctl` command
 - `httpd`
 - Command that can be used to test config files of virtual hosts
 - `httpd -t -D DUMP_VHOSTS`
 - Displays all configured virtual hosts

Apache

- Virtual Hosts
 - Run multiple sites from one server
 - Name-based
 - Uses the domain name in the request to forward to the correct virtual host
 - IP-based
 - Uses the IP-address in the request to forward to the correct virtual host
 - Name-based are far more common than IP-based
 - Often configured in separate configuration files inside `/etc/httpd/conf.d/*.conf`

Apache

- Virtual Hosts

- Name-based virtual host example

```
<VirtualHost *:80>
    DocumentRoot "/www/example1"
    ServerName www.example.com

    # Other directives here
</VirtualHost>

<VirtualHost *:80>
    DocumentRoot "/www/example2"
    ServerName www.example.org

    # Other directives here
</VirtualHost>
```

Apache

- Virtual Hosts
 - Name-based virtual host example with ports

```
<VirtualHost 172.20.30.40:80>  
    ServerName www.example.com  
    DocumentRoot "/www/domain-80"  
</VirtualHost>
```

```
<VirtualHost 172.20.30.40:8080>  
    ServerName www.example.com  
    DocumentRoot "/www/domain-8080"  
</VirtualHost>
```

Apache

- Apache and SELinux
 - Security Enhanced Linux
 - Additional layer of security on many Linux distributions
 - Requires extra configuration for Apache
- Service on different port
 - `sudo semanage port -a -t http_port_t -p tcp [port_number]`
 - `sudo semanage port -l`
 - List all ports that are usable through selinux
- Allow different content folders
 - `sudo setsebool -P httpd_unified 1`

Apache

- Modules



**Add functionality
to your server**



**Many modules
available**



**Load only what's
needed**



**Write your own
modules**

Apache

- Modules
 - Install through the package manager of your Linux distribution.
 - List of all modules that come with Apache:
 - httpd.apache.org/docs/current/mod/
 - Location of modules:
 - `/etc/httpd/modules`
 - Symlink to external path for easier configuration with relative paths towards `ServerRoot`
 - Location of configuration files:
 - `/etc/httpd/conf.modules.d/`

Apache

- Modules

- Syntax for loading modules:

- `LoadModule [Module_name] [path_to_module_binary]`
 - e.g. `LoadModule alias_module modules/alias.so`

- Display list of all loaded modules:

- `httpd -D DUMP_MODULES`

Apache

- Modules
 - Frequently used Apache modules

Module Name	Purpose
<code>mod_ssl</code>	Implements SSL and TLS
<code>mod_alias</code>	Simple URL remapping
<code>mod_rewrite</code>	Rule based remapping of URLs
<code>mod_status</code>	Provides information on server activity and performance
<code>mod_deflate</code>	Compression before content is delivered to the client
<code>mod_cgi</code>	Application execution as defined in the script
<code>mod_cgid</code>	Application execution using an external CGI daemon

Webservers overview

- NginX
 - Since 2004
 - Open-source
 - Free
 - Cross-platform: Linux, Unix, BSD, but also Windows
 - Vast majority on Linux
 - Initial design aimed towards webserver and proxy

Webservers overview

- NginX
 - Characteristics
 - Light on resources
 - Great under heavy load
 - Extensible through modules, but often more complex configuration
 - Great for delivering static content
 - (Reverse) proxy
 - Load balancer
 - No directory based configuration
 - Faster
 - Centralized control of configuration is often more secure
 - Web request are handled mainly based on the URI

NginX

- Configuration
 - The NginX configuration was largely based on the Apache configuration
 - Many directives from Apache with the same name in NginX
 - Admins with knowledge in Apache adapt easily to NginX

NginX

- Configuration files
 - NginX wide main configuration file:
 - `/etc/nginx/nginx.conf`
 - Additional configuration files:
 - Extra configuration files can be added and will be read by Nginx if placed in:
 - `/etc/nginx/conf.d/*.conf`
 - Module configuration files:
 - `/usr/share/nginx/modules/*.conf`

NginX

- Directives
 - Rules, instructions and settings for how the webserver should behave.
 - Simple directives:
 - name and parameters separated by spaces and ends with a semicolon (;)
 - Blocked directives:
 - same structure as a simple directive
 - instead of the semicolon it ends with a set of additional instructions surrounded by curly brackets ({ })
 - can have other directives inside braces: context
 - Directives placed in the configuration file outside of any contexts
 - Considered to be in the main context
 - `events` and `http` directives: main context
 - `server` in `http`
 - `location` in `server`

NginX

- Commonly used directives
 - `include`
 - Which config files to include
 - e.g. `include /etc/nginx/conf.d/*.conf;`
 - `user`
 - Which user and which group to run NginX as
 - Default `nginx` in CentOS

NginX

- Commonly used directives
 - `http`
 - Main context for the HTTP server directive
 - `server`
 - Used to define a `server` block (see Server Blocks)
 - `location`
 - Used to define a `location` block (see Location Blocks)
 - Located within a `server` block

NginX

- Commonly used directives

- `listen`

- Sets the address and port for IP, or the path for a UNIX-domain socket on which the server will accept requests

- Default port 80

- Default all IP addresses of the server

- `Listen 80;` `#Listen on all IPv4-addresses on port 80`
 - `Listen [::]:80;` `#Listen on all IPv6-addresses on port 80`
 - `Listen 10.3.50.50:80` `#Listen on Listen on IPv4-address 10.3.50.50 on port 80`

- Many other configuration options for listen directive available

NginX

- Commonly used directives

- `root`
 - Path to the folder that contains the actual web documents
- `server_name`
 - Used to define the domain name that will answer the incoming request

NginX

- NginX commands
 - `sudo systemctl`
 - `start nginx`
 - Starts the NginX service
 - `status nginx`
 - Command to check the status of the NginX service
 - Running (active (running))
 - Start with the Operating System (enabled)

NginX

- NginX commands
 - `nginx`
 - The main command used for managing NginX
 - `-s`
 - `stop` – fast shutdown
 - `quit` – graceful shutdown
 - `reload` – reloading the configuration file
 - `reopen` – reopening the log files
 - `-t`
 - Test the configuration file(s)
 - `-v`
 - Display the installed version of NginX

NginX

- Server Blocks

- Configuration blocks containing separated configuration contexts
- Hierarchically structured
- Every request runs through the hierarchy to determine which configuration block is applied
 - First selected on the `listen` directive, then the `server_name` directive
 - In case of competing directives, the one that is more specific will win
 - e.g. `listen 192.168.1.10;` will win over `listen 80;`, even if the second one has a nameserver defined, while the first one doesn't.

NginX

- Server Blocks

- `server` block

- Subset of configurations that define a virtual server
 - Multiple server blocks are possible to decide which block will handle the request based on domain name, IP address and port
 - Similar to Virtual Hosts in Apache: run multiple sites from one server

- `location` block

- Located within a server block
 - Defines how requests are processed for different URIs and resources

NginX

- Server Block examples
 - Name-based Server Blocks, serving static files

```
server {  
    server_name www.domain1.com;  
    access_log logs/domain1.access.log main;  
  
    root /var/www/domain1.com/htdocs;  
}  
  
server {  
    server_name www.domain2.com;  
    access_log logs/domain2.access.log main;  
  
    root /var/www/domain2.com/htdocs;  
}
```

NginX

- Server Block examples
 - Port-based Server Blocks, serving static files

```
server {  
    listen 80;  
    server_name www.domain.com;  
    access_log logs/domain.access-80.log main;  
  
    root /var/www/domain.com-80/htdocs;  
}  
  
server {  
    listen 8080;  
    server_name www.domain.com;  
    access_log logs/domain.access-8080.log main;  
  
    root /var/www/domain.com-8080/htdocs;  
}
```