

Containerization

Server OS / Operating Systems II

Contents

- Container architecture
- Microservices and Cloud-Native
- Orchestration
- Docker concepts
- Docker installation
- Docker commands
- Docker Compose

Container architecture

- **Applications/services:** crucial in the operation of organizations
- How to provision them
 - **Securely**
 - **Efficiently**


=> isolated from each other

=> only use the resources it needs

=> Scalable


Container architecture

- **Before 2000:**
 - Very often **1 application/service for each hardware server**
 - hardware very often overprovisioned
 - => Waste of resources
 - Hardware
 - Support
 - Uptime



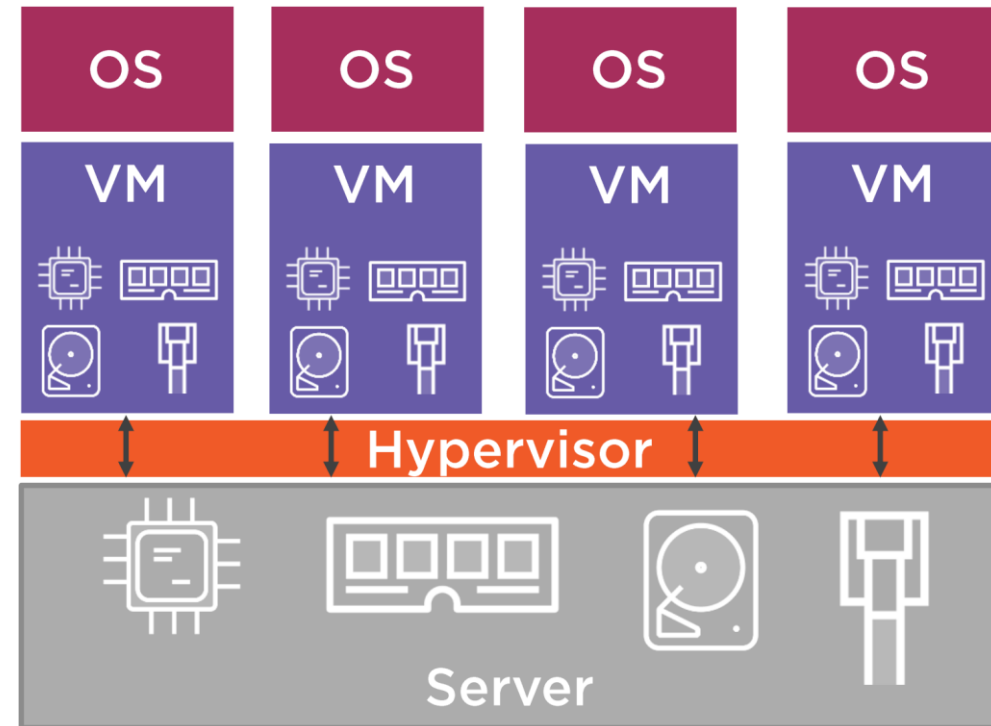
Container architecture

- **From 2000 onwards:**
 - **Virtualization** gets popular
 - Hardware:
 - shared between several Virtual Machines (VM's)
 - Less waste of resources



Container architecture


- **From 2000 onwards:**
 - **Every app:**
 - **its own VM**
 - its own virtual hardware
 - Its own virtual Operating System
 - Which requires resources for every VM
 - Often the same full OS in different VM's



Container architecture

- **Containers**

- **1 app per container**
- Containers are isolated
- each container **separate namespaces**:
 - filesystem
 - process tree
 - users and groups tree
 - network stack



Container architecture

- **Containers**

- act like a complete separate OS's
- **share the same OS kernel with the host computer**
 - => Only one OS kernel necessary
 - => less resources needed

Container architecture

- **Containers**

- => **container OS kernel = host OS kernel**
 - Linux containers on Linux hosts
 - Windows containers on Windows hosts
 - VM's can be used as host for mixed solutions
 - Linux container on Windows host with a Linux VM in Hyper-V or WSL

Container architecture

- Containers have existed in the Linux world for decades
 - e.g. Google's search engine
 - Complicated



- **Containers made easy:**
cgroups and namespaces into manageable containers

Container architecture




- Controls individual containers
 - start and stop
- **Standard in containerization**
 - **Mainly in Linux**
 - also for Windows and Mac
 - Platform independent:
 - cloud, on-prem, hybrid, Linux, Mac, Windows
 - commands are the same

Microservices and Cloud-Native

- **Monolithic app architecture**

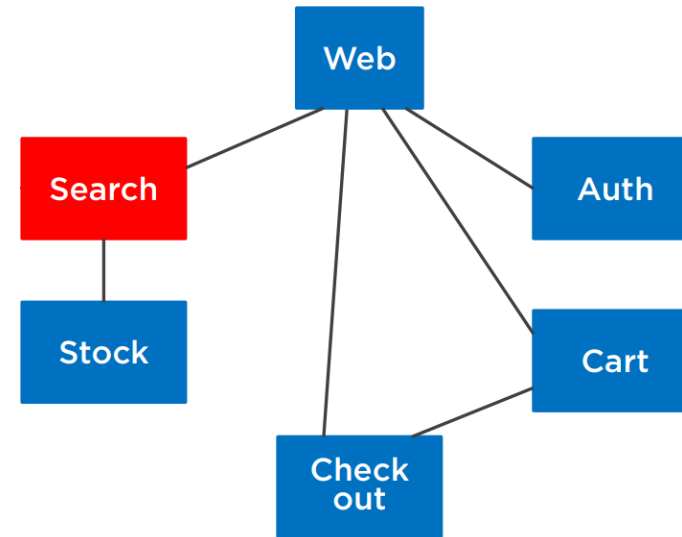
- Large applications
- all functionality integrated into one big binary
- fix/upgrade of one component (e.g. search)?
=> Everything down!



Microservices and Cloud-Native

- **Microservices:**

- All the separate components of an application => **split up**
- fix/upgrade of one component (e.g. search)?
=> Only that component down!




Microservices and Cloud-Native

- Containers are used as microservices
- **One app/process per container**
 - Scalable
 - Self-healing
 - Portable
 - Efficient in resource usage

Microservices and Cloud-Native

- **Docker itself:** written in a Microservices architecture
 - **Split up in different components:**
 - **Daemon (dockerd)**
 - Provides API for the client
 - Manages the actual containers
 - **CLI Client (docker)**
 - Allows interacting with the containers
 - e.g. start, stop
 - Connects to the daemon (local or remote) for the operations
 - Can be upgraded without stopping containers




Microservices and Cloud-Native

- **Cloud-Native apps:**
 - microservices optimized for the cloud
 - open source
 - containerized
 - dynamically orchestrated
 - microservices-oriented

Orchestration

- Microservices architecture can become **complex**...
 - Single app per container
=> often requires **many containers for the full stack**
 - Often multiple instances of the same container
 - Often complex communication layout between containers



Orchestration

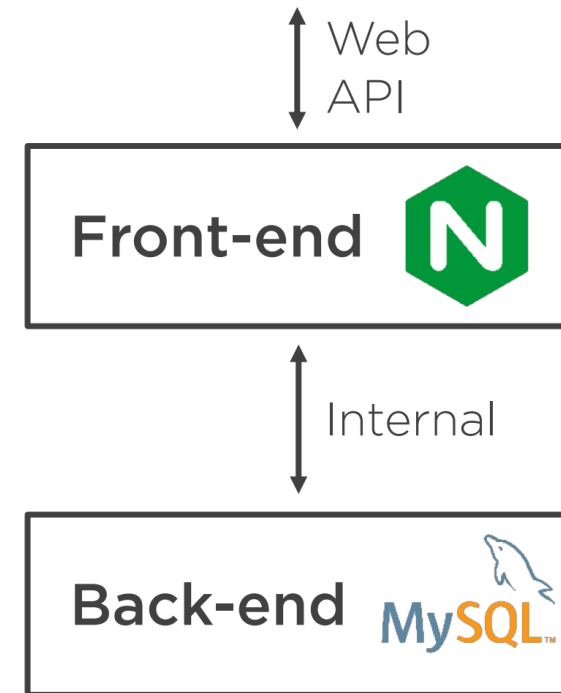
- Microservices architecture can become **complex**:
- Containers
 - **dynamically** managed (according to load and failures)
 - often spread-out over **multiple hosts**
 - On-premises (on-prem), cloud or hybrid

Orchestration

- Orchestration: **automation of management of container stacks**
- cfr orchestra:
 - every musician = container
 - conductor of the orchestra
 - starts and stops different groups
 - sets the tempo
 - manages the stack


Orchestration

- E.g. webstack:
- NginX as web front-end
- MySQL as backend
- Under normal load (desired state):
 - 2 instances of NginX container
 - (for load balancing)




Orchestration

- If the 2 existing instances of NginX (3 and 5) get too much load:
 - 2 more instances (2 and 4)
 - Automatically added by the orchestrator




Orchestration

- Load = automatically balanced among all instances




Orchestration

- If load is back to normal:
 - excessive instances removed
 - until the desired state (2 NginX, 1 MySQL) is reached.




Orchestration

- If instance fails:
 - immediately picked up by the orchestrator
 - replacement instance started up



Orchestration

- **Docker Swarm:** Docker's own orchestration
 - Easy to use
 - All basic functionality included



Orchestration

- **Kubernetes (K8s):** Google's orchestration
 - **Large amount of functionality**
 - Becoming the de facto **standard**
 - Integrated in all major Cloud Services
 - Integrated in many Server Class Operating Systems
 - Integrated in Docker

[Kubernetes \(K8s\)](https://kubernetes.io/) is an open-source system for automating deployment, scaling, and management of containerized applications.



Docker concepts

Docker Images

- Basis for every container
- **'template'** for a container
- Read-only -> **immutable**
- **Build-time construct**
- “stopped” container
- Can be pulled from a registry with `docker pull` command

Docker concepts

Docker Images

- Contains:
 - **several stacked layers**
 - to build a unified filesystem
 - json manifest file
 - describes the image
 - how the layers should be combined together
 - e.g.
 - bottom layer for the OS files
 - next layer for the app files
 - third layer for updates

Docker concepts

Docker Images

- Layers are **locally stored**
 - in Linux: /var/lib/docker/[storage_driver_name]/diff
 - In Windows: C:\ProgramData\Docker\Windows Filter
- **extra writable layer**
 - added when a container is created from the image

Docker concepts

Docker Images

- Are stored in **registries**
- Need to have a copy in the **local registry** on the host
- If a local copy is not available
=> **automatically downloaded** ('pulled') from an image registry
to the local registry

Docker concepts

Docker Images

- Docker registries

- Docker hub** (hub.docker.com)

- Default
 - Contains thousands of images for applications

Docker concepts

Docker Images

- Docker registries

- Docker hub** (hub.docker.com)

- Official images:

- maintained by the developer of the app
 - Should be: stable, up-to-date, tested and well documented
 - Don't need a separate namespace
 - e.g. `nginx`

Docker concepts

Docker Images

- Docker registries

- Docker hub** (hub.docker.com)

- Unofficial images: not from the official developer of the app
 - Require a separate namespace
 - e.g. `nginxdemos/hello`
 - Many different versions
 - `latest` version: most up-to-date stable version (usually)

Docker concepts

Docker Images

- Docker registries
 - Other public registries
 - Google
 - Amazon
 - Microsoft
 - ...
 - Private registries
 - provided and maintained by your own organisation
 - privately created images

Docker concepts

Docker Images

- Image naming syntax: **registry/repository:tag**
e.g. `docker.io/ubuntu:latest`

registry: name of the registry

- Default: `docker.io` (docker hub)
- If default
=> does not need to be mentioned

Docker concepts

Docker Images

- Image naming syntax: **registry/repository:tag**
e.g. `docker.io/ubuntu:latest`

repository: separate space in a registry

Docker concepts

Docker Images

- Image naming syntax: **registry/repository:tag**
e.g. `docker.io/ubuntu:latest`

tag: name of the image in the repository

- Default: `latest`
- If default
=> does not need to be mentioned
- `latest` is tagged as such manually by the repository maintainer

Docker concepts

Docker Images

- Are built for a **specific kernel**
 - Windows images are a lot bigger
 - necessary for apps that need a Windows kernel
 - e.g. Docker images for Powershell

Image	SIZE
Windows image for Powershell	5.35GB
Linux image for Powershell	339MB

Docker concepts

Docker Images

- Base images for building app image are focused on being very **lightweight**
 - Linux: **alpine**
 - Windows: **nano server**

Docker concepts

Docker containers

- Based on image
- **Runtime construct:** “running” instance of image
- Multiple instances of same image possible

Docker concepts

Docker containers

- Should be **ephemeral**:
 - only used for set period of time
 - changes need to be made in the image
 - replace container with a new one based on the new image
- **Stopping a container**:
 - the container is not removed but exited
 - Data persists in the container

Docker concepts

Docker containers

- Created to be lightweight
- Usually only one process running
 - ...in Linux containers
 - The Windows kernel needs more processes
=> Windows containers too

Docker concepts

Docker containers

- Created to be lightweight

Example:

```
docker container run -it ubuntu bash
```

=> interactive ubuntu docker

```
vi
```

=> will not work

vi is not available in this image (lightweight)

```
top
```

=> only bash and top (forked from bash) are running

Docker concepts

Docker containers

- Created to be lightweight
 - Example:
exit the terminal
 - => the bash process is only running process
 - => bash will be stopped
 - => the container itself will also stop
 - To exit an interactive container without stopping
 - => do [CTRL+P+Q]

Docker concepts

Docker containers

Processtable of a Powershell Windows Container

PS C:\> ps


NPM(K)	PM(M)	WS(M)	CPU(s)	Id	SI	ProcessName
6	1.02	4.34	0.00	1168	1	CExecSvc
5	0.91	2.19	0.00	1728	1	CompatTelRunner
7	1.17	4.81	0.03	1244	1	conhost
10	6.72	8.16	0.03	1736	1	conhost
11	1.79	4.79	0.19	984	1	csrss
5	0.82	2.62	0.02	616	1	fontdrvhost
0	0.06	0.01	0.00	0	0	Idle
22	4.21	12.37	0.14	520	1	lsass
58	43.09	78.44	2.09	1264	1	pwsh
11	2.17	5.96	0.08	500	1	services
4	1.19	1.24	0.31	956	0	smss
13	2.87	9.98	0.05	624	1	svchost
16	2.41	7.96	0.08	1036	1	svchost
25	6.09	18.20	0.11	1096	1	svchost
15	3.04	8.76	0.05	1144	1	svchost
18	7.53	13.60	0.09	1224	1	svchost
34	5.45	16.94	1.20	1344	1	svchost
22	11.73	24.00	0.78	1456	1	svchost
7	1.34	5.60	0.02	1468	1	svchost
8	1.56	6.17	0.03	1496	1	svchost
0	0.16	0.13	0.66	4	0	System
12	1.79	6.93	0.08	284	1	wininit

Processtable of a Powershell Linux Container

PS /> ps

PID	TTY	TIME	CMD
1	pts/0	00:00:00	pwsh
39	pts/0	00:00:00	ps

- Docker concepts
Docker images vs containers



Docker concepts

Building Docker images

- Similar to building an app outside without container
- Steps are defined in a **Dockerfile**
 - FROM:
base image for new image
 - WORKDIR:
directory in your image filesystem
all actions should be taken here

Docker concepts

Building Docker images

- COPY
copies files from the host to the image
e.g. your application files
- EXPOSE
documents which ports the application uses
- RUN
run a command
e.g. to build the application

Docker concepts

Building Docker images

- Additional metadata can be added to the Dockerfile
 - how to run the container based on this image
 - CMD
default process to run in the container

Docker concepts

Building Docker images

- The actual image can be built from the `Dockerfile` with the command...
 - `docker image build -t [image_name:tag_name] [path_to_dir_with_all_necessary_files]`
 - e.g. `docker image build -t testapplication:1.0 /home/user/testapp`
- This image can be used to start a container

Docker concepts

Building Docker images

- Building small images
 - Use **small base images** (cfr FROM)
 - Use **multi-stage builds**
 - The first stage
builds the app with build tools in a temporary image
 - Second stage
copies the built app over to the final image without the build tools

Docker concepts

Persisting data: volumes

- More info: <https://docs.docker.com/storage/volumes/>
- Containers should be
ephemeral
- **Where to keep data** after removing the container?
=> Volumes

Docker concepts

Persisting data: volumes

- **Volumes**

- allow **mapping folders** from inside the container to the host
- container is removed:
 data remains
- can be shared among different containers
- Can be on a remote host, cloud storage, SAN, NAS...
- Will by default be created on the host:
 - Linux: `/var/lib/docker/volumes`

Docker concepts

Connecting to containers: port mapping

- The network stack of a container is isolated from the host
- How to access the container from the outside/host?
=> **Port mapping**

Docker concepts

Connecting to containers: port mapping

- Map a port from the host port to a container port
 - **Host port:**
 - Host will listen to traffic coming in on this port and forward it to the...
 - **Container port:**
 - the container will accept traffic on this port.
 - Should be the same port as the one stated with EXPOSE during build.
- Port mapping syntax:
 - [host_port]:[container_port]
 - e.g.: 80:80

Docker installation

- All the different procedures for installing Docker in different environments are available at <https://docs.docker.com>
- Docker Desktop
 - Ideal for testing and development
 - Windows
 - Requires the Containers and Hyper-V features
 - VirtualBox will not work anymore
 - Also possible through Chocolatey
 - Can run from WSL2 too for linux container support
 - MAC OS

Docker installation

- Docker in Windows Server
 - Ideal for production environments
 - Through Powershell

```
Install-Module -Name DockerMsftProvider -Repository PSGallery  
-Force
```

```
Install-Package -Name docker -ProviderName DockerMsftProvider
```

Docker installation

- Docker in Linux
 - Suitable for
 - production
 - development
 - testing environments
 - Procedure depends on Linux distribution
 - <https://docs.docker.com>
 - If available, install using the repositories

Docker commands

`docker`

- Displays a quick overview of available docker options and parameters

`docker version`

- Displays the installed version of docker

`docker info`

- Displays general information of your docker installation

Docker commands

```
docker [image] pull image_name
```

- Pulls the latest image from the registry
- Based on image_name

```
docker [image] pull image_name[:tag]
```

- Pulls an image from the registry
- Based on image_name
- Optional tag referring to version of the image

Docker commands

```
docker [container] run image_name
```

- Runs a container
- Based on `image_name`
- If the image is not present:
=> image downloaded (pulled) from the registry

Docker commands

```
docker [container] run [options] image_name
```

OPTIONS

- **-d : detached**
container runs in the background
- **-i -t : interactive and TTY -pseudoterminal**
gives a cli terminal into the container
- **-p port_host:port_container :**
maps the internal port_container to the port_host
makes internal port available from the host

Docker commands

```
docker [container] run [options] image_name
```

OPTIONS

- **-e environment variable**
sets environment variable
- **--name container_name**
names the newly started container with name `container_name`
- **--mount source=volume_name,target=path_to_targetfolder**
- **-v path_to_hostfolder:path_to_containerfolder**
mounts the volume on the host with name `volume_name`
to the folder inside the container `path_to_target_folder`

Docker commands

`docker image`

- Manage images
- **`ls` : list**
List all the images available locally on the host
Same as `docker images` command
- **`rm image_name` : remove**
Remove image with name `image_name`
Same as `docker rmi image_name`

Docker commands

`docker image`

- **`build -t image_name:image_tag path_to_build_files:`**
build a new image with name **`image_name`** and tag **`image_tag`**
needs `path_to_build_files` pointing to the build files, including the Dockerfile
- **`prune`**
Remove all dangling images
- **`prune -a`**
Remove all unused images

Docker commands

`docker image`

- `$(docker images -q)` = all images

Can be used with `rm`

e.g. `docker image rm $(docker images -q)`
removes all images

Docker commands

`docker container`

- Manage containers
- **`ls` : list**
List all the running containers on the host.
Same as `docker ps` command
- **`ls -a` : list all**
List all the containers on the host.

Docker commands

`docker container`

- **`rm container_name/id:remove`**

Remove container with name or id `container_name/id`

Same as `docker rm container_name/id`

- **`exec command :execute`**

Run the command `command` inside the container (must be running)

Can be used with the option `-it` to run the command `command` interactively in the container

Docker commands

`docker container`

- **`logs container_name/id`**
shows the logs for the container with name or id `container_name/id`
- **`top container_name/id`**
shows the running processes int the container with name or id `container_name/id`
same as `top` in a regular linux environment
- **`prune`**
removes all stopped containers

Docker commands

`docker volume`

- **`create volume_name`**
creates a volume with name `volume_name`
- **`ls`**
list all volumes on the host
- **`rm volume_name`**
removes volume with name `volume_name`

Docker commands

`docker system`

- **prune** : removes all...
 - all stopped containers
 - all networks not used by at least one container
 - all dangling images
 - all dangling build cache
- **prune -a** : removes all...
 - all stopped containers
 - all networks not used by at least one container
 - all images without at least one container associated to them
 - all build cache

Docker-Compose

Working with commands:

- Practical for individual containers
- Less so for a stack of containers with plenty of additional parameters

⇒ **Docker-Compose**

- Tool for defining and running multi-container Docker applications
- Defined in 1 yaml file
- Needs to be installed: <https://docs.docker.com/compose/install/>

Docker-Compose

docker-compose.yaml example:

services:

- traefik:
- image: traefik:alpine
- container_name: traefik
- command: --api --docker
- volumes:
- - /var/run/docker.sock:/var/run/docker.sock
- - /dl/config/traefik/traefik.toml:/etc/traefik/traefik.toml
- ports:
- - "80:80"
- environment:
- - PGID=1000
- - TZ=Europe/Brussels
- restart: always

Docker-Compose

Start:

V1: docker-compose up -d

V2: docker compose up -d

Stop:

V1: docker-compose down -d

V2: docker compose up -d