

Samenvatting theorie Containerization

Server OS

Bachelor Toegepaste Informatica
Graduaat Netwerken

Auteur : Geert Coulommier

Design, Technologie en Gezondheidszorg

1. Contents


- Container architecture
- Microservices and Cloud-Native
- Orchestration
- Docker concepts
- Docker installation
- Docker commands

2. Container architecture

- Applications / services are crucial in the operation of organizations
- How to provision them in a secure and efficient way?
 - Should be isolated from each other
 - Should only use the resources it needs
 - Scalable


2.1 Before 2000:

- Very often 1 application / service for each hardware server
- To be on the safe side for future upscaling of the application the hardware was very often overprovisioned
- Waste of resources
 - Hardware
 - Support
 - Uptime




2.2 From 2000 onwards:

- Virtualization gets popular
- Hardware is shared between several Virtual Machines (VM's) through the hypervisor
- Less waste of resources



Every app gets its own VM with its own virtual hardware and virtual Operating System


- Which requires resources for every VM
- Often the same OS in different VM's



<https://app.pluralsight.com/library/courses/docker-getting-started/exercise-files>

2.3 Containers

- Apps are isolated from one another in separated containers, with for each container separate namespaces:
 - filesystem
 - process tree
 - users and groups tree
 - network stack



<https://app.pluralsight.com/library/courses/docker-getting-started/exercise-files>

- Each container acts like a complete separate OS
- Resources used by containers are limited by control groups (cgroups)
- Containers share the same OS kernel from the host computer
- Only one OS kernel necessary, so less resources needed
- As containers use the kernel for the host, containers need the same kernel as the host
 - Linux containers on Linux hosts
 - Windows containers on Windows hosts
 - VM's can be used as host for mixed solutions
 - Linux container on Windows host with a Linux VM in Hyper-V
- One application per container makes them
 - Scalable
 - Self-healing
 - Portable
 - Lightweight and fast
 - Efficient in resource usage
- Containers have existed in the Linux world for decades

- e.g. Google's search engine




2.4

- Docker made the controlling of namespaces and cgroups easy by bundling them in easy manageable containers
- Controls individual containers, in essence starting and stopping
- De facto standard in containerization
 - Mainly in Linux, but also for Windows and Mac
 - Platform independent:
 - cloud, on-prem, hybrid, Linux, Mac, Windows: commands are the same

3. Microservices and Cloud-Native

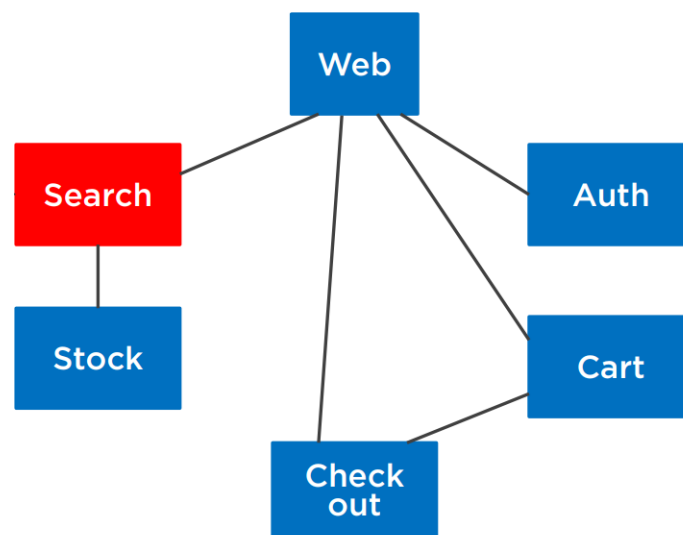
- Monolithic app architecture
 - Large applications with all functionality integrated in one big binary
 - The whole entity had to be taken down to fix/upgrade a single item (e.g. search)



<https://app.pluralsight.com/library/courses/docker-kubernetes-big-picture/exercise-files>


3.1 Microservices:

- All the separate components of an application are split up in separate components
- When one component needs to get fixed/updated, only that component needs to be taken down and replaced



<https://app.pluralsight.com/library/courses/docker-kubernetes-big-picture/exercise-files>

- Containers are used as microservices
- One app/process per container
 - Scalable
 - Self-healing
 - Portable
 - Efficient in resource usage
- Docker itself has also been written in a Microservices architecture
 - Split up in different components that can be managed and changed separately from one another
 - Client-Server model:
 - CLI Client (docker)
 - Allows interacting with the containers, like starting and stopping commands
 - Connects to the daemon (local or remote) for actual operations
 - Daemon (dockerd)
 - Provides API for the client to interact with
 - Manages the actual containers
- This architecture allows to upgrade the client without having to take down running containers




3.2 Cloud-Native

- Optimize microservices with containers for the cloud and you get Cloud-Native apps
- “Cloud native computing uses an open source software stack to be:
 - Containerized. Each part (applications, processes, etc) is packaged in its own container. This facilitates reproducibility, transparency, and resource isolation.
 - Dynamically orchestrated. Containers are actively scheduled and managed to optimize resource utilization.
 - Microservices-oriented. Applications are segmented into microservices. This significantly increases the overall agility and maintainability of applications.”

<https://www.cncf.io/about/faq/>

4. Orchestration

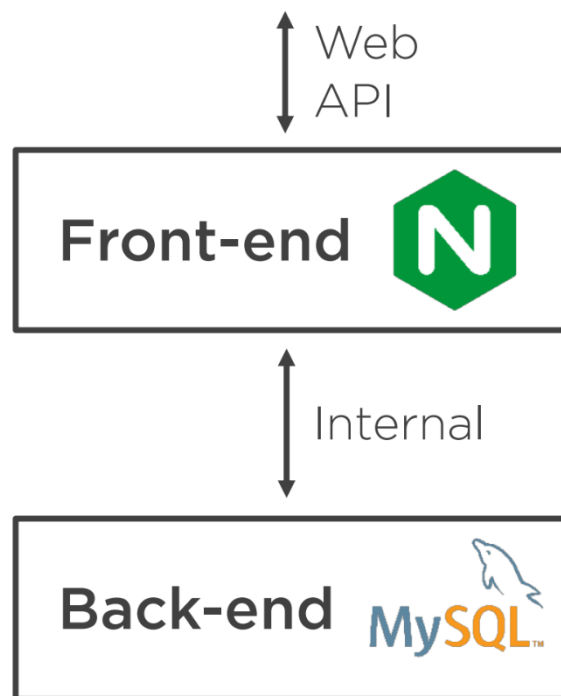
- Microservices architecture can become complex...
 - Single app per container often requires many containers for the full stack
 - Often multiple instances of the same container
 - Containers need to communicate with one-another in often complex layout



<https://app.pluralsight.com/library/courses/docker-kubernetes-big-picture/exercise-files>


- But also... microservices architecture can become complex
 - Containers need to be dynamically managed according to load and failures
 - Containers are often spread-out over multiple hosts
 - On-premises (on-prem), cloud or hybrid
- Orchestration will automate the task related to the management of container stacks
 - Comparison: orchestra: every musician is like a container, where the conductor of the orchestra starts and stops different groups, sets the tempo, and as such manages the stack

- For instance, consider a stack with NginX as web front-end and MySQL as backend. For load balancing, 2 instances of the NginX container are desired under normal load.




<https://app.pluralsight.com/library/courses/docker-kubernetes-big-picture/exercise-files>

- The orchestrator will automatically add 2 more instances of NginX (2 and 4) if the 2 existing ones (3 and 5) get too much load.




<https://app.pluralsight.com/library/courses/docker-kubernetes-big-picture/exercise-files>

- The load will automatically be balanced among all instances.




<https://app.pluralsight.com/library/courses/docker-kubernetes-big-picture/exercise-files>

- And when the load is back to normal, excessive instances will be removed until the desired state (2 NginX, 1 MySQL) is reached.



<https://app.pluralsight.com/library/courses/docker-kubernetes-big-picture/exercise-files>

- Also, if an instance were to fail, it will immediately be picked up by the orchestrator, and a replacement instance will be started up.



<https://app.pluralsight.com/library/courses/docker-kubernetes-big-picture/exercise-files>

- Docker Swarm: Docker's own orchestration
 - Easy to use
 - All basic functionality included
- Kubernetes (K8s): Google's orchestration
 - Large amount of functionality
 - Becoming the de facto standard
 - Integrated in all major Cloud Services
 - Integrated in many Server Class Operating Systems
 - Integrated in Docker

5. Docker concepts

5.1 Docker Images

- Basis for every container, 'template' for a container
- Read-only -> immutable
- Build-time construct
- "stopped" container
- Is run with the `docker run` command to start a container
- Can also be pulled from a registry without immediately starting a container with the `docker pull` command
- Contains several stacked layers to build a unified filesystem, and a json manifest file, describing the image and how the layers should be combined together
 - e.g. bottom layer for the OS files, next layer for the app files, third layer for updates
 - Layers are locally stored
 - in Linux: `/var/lib/docker/[storage_driver_name]/diff`
 - In Windows: `C:\ProgramData\Docker\Windows Filter`
 - An extra writable layer is added when a container is created from the image
- Are stored in registries
- Needs to have a copy in the local registry on the host
- If a local copy is not available, it will be automatically downloaded ('pulled') from an image registry to the local registry

5.1.1 Docker registries

- Docker hub (hub.docker.com)
 - Default
 - Contains thousands of images for applications
 - Official images: maintained by the developer of the app
 - Should be stable, up-to-date, tested and well documented
 - Don't need a separate namespace
 - e.g. `nginx`
 - Unofficial images: not from the official developer of the app
 - Require a separate namespace
 - e.g. `nginxdemos/hello`
 - Many provide different versions, with the latest version usually being the most up-to-date stable version
- Other public registries, like Google, Amazon, Microsoft...
 - Private registries, e.g. provided and maintained by your own organisation with privately created images

5.1.2 Image naming syntax: **registry/repository:tag**

- e.g. `docker.io/ubuntu:latest`
- **registry**: name of the registry
 - Default: `docker.io` (docker hub)
 - If the registry is docker hub (`docker.io`) it does not need to be mentioned (default)
- **repository**: separate space in a registry
- **tag**: name of the actual image in the repository
 - If the tag is `latest`, it does not need to be mentioned (default)
 - `latest` is tagged as such manually by the repository maintainer
 - Usually refers to latest stable
 - Can be older than other available images, like `edge` (development unstable version)

5.1.3 Linux or Windows Images

- Images are built for a specific kernel
- Windows images are a lot bigger, but necessary for apps that need a Windows kernel

Image	SIZE
Windows image for Powershell	5.35GB
Linux image for Powershell	339MB


- Base images on which to build an app image are focused on being very lightweight:
 - Linux: `alpine`
 - Windows: `nano server`

5.2 Docker containers

- Based on image
- Runtime construct: "running" instance of image
- Multiple instances of same image possible
- Are meant to be ephemeral: only used for set period of time. If changes need to be made, they are done in the image, the container is stopped and replaced with a new container based on the new image.
- When stopping a container, the container is not removed but exited
 - Data persists in the container

- Created to be lightweight
 - Usually only one process running
 - ..in Linux containers. The Windows kernel needs more processes, therefore Windows containers too.
 - e.g. Try `docker run -it ubuntu bash` for an interactive ubuntu docker and then try `vi`, it will not work because `vi` is not available in this image (lightweight). Run `top` to see only `bash` and `top` (forked from `bash`) running.
 - When exiting the terminal, the `bash` process will be stopped, the only running process in the container. Therefore, the container itself will also stop.
 - To exit an interactive container without stopping, do `[CTRL+P+Q]`

5.3 Docker images vs containers



<https://app.pluralsight.com/library/courses/docker-deep-dive-update/exercise-files>

5.4 Building Docker images

- Steps are very similar to building an app outside of a container environment
- Steps to create an image are defined in a `Dockerfile`
 - **FROM:** specifies the image this new image should be based on
 - **WORKDIR:** specifies the directory in your image filesystem in which all subsequent actions should be taken
 - **COPY:** copies files from the host to the image, e.g. your application files
 - **EXPOSE:** documents which ports the application uses
 - **RUN:** run a command, e.g. to build the application based upon the imagefiles copied in the previous `COPY` parameter
- Additional metadata about how to run the container based on this image can also be added to the `Dockerfile`
 - **CMD:** default process to run in the container
- The actual image can be built from the `Dockerfile` with the command...
- `docker image build -t [image_name:tag_name] [path_to_dir_with_all_necessary_files]`
- e.g. `docker image build -t testapplication:1.0 /home/user/testapp`
- This image can be used to start a container
- Building small images
 - Use small base images (`FROM`)
 - Use multi-stage builds

- The first stage builds the app with the build tools in a temporary image
- Second stage copies the built app over to the final image without the build tools

• Docker concepts

5.5 Persisting data: volumes

- More info: <https://docs.docker.com/storage/volumes/>
- Containers are meant to be ephemeral
- Where to keep data after stopping/removing the container?
- Volumes
 - allow mapping folders from inside the container to the host
 - container is stopped or removed: data remains
 - can be shared among different containers
 - Can be on a remote host, cloud storage, SAN, NAS...
 - Will by default be created on the host:
 - Linux: /var/lib/docker/volumes

6. Docker installation

- All the different procedures for installing Docker in different environments are available at <https://docs.docker.com>
- **Docker Desktop**
 - Ideal for testing and development
 - Windows
 - Requires the Containers and Hyper-V features
 - VirtualBox will not work anymore
 - Also possible through Chocolatey
 - Can run from WSL2 too to for linux container support
 - MAC OS
- **Docker in Windows Server**
 - Ideal for production environments
 - Through Powershell
 - Install-Module -Name DockerMsftProvider -Repository PSGallery -Force
 - Install-Package -Name docker -ProviderName DockerMsftProvider
- **Docker in Linux**
 - Suitable for production, development and testing environments
 - Procedure depends on Linux distribution
 - <https://docs.docker.com>
 - If available, install using the repositories

7. Docker commands

`docker`

- Displays a quick overview of available docker options and parameters.

`docker version`

- Displays the installed version of docker.

`docker info`

- Displays general information of your docker installation.

`docker [image] pull (image_name)`

- Pulls the latest an image from the registry based on (image_name).

`docker [image] pull (image_name):(tag)`

- Pulls an image from the registry based on (image_name), with an optional tag referring to version of the image

`docker [container] run (image_name)`

- Runs a container based on (image_name). If the image for the container is not present, it will be downloaded (pulled) from the registry.
- `-d` option: detached: container runs in the background
- `-i -t` option: interactive and TTY -pseudoterminal: gives a cli TTY into the container
- `-p (port_host:port_container)` option: maps the internal (port_container) to the (port_host), making it available from the host
- `-e (environment variable)` option: sets environment variable
- `--name (container_name)` option: names the newly started container with name (container_name)
- `--mount source=(volume_name),target=(path_to_targetfolder)` option: mounts the volume on the host with name (_volume_name) to the folder inside the container (path_to_target_folder)

`docker image`

- Manage images
- `ls` command: list. List all the images available locally on the host
 - Same as `docker images` command
- `rm (image_name)` command: remove. Remove image with name image_name.
 - Same as `docker rmi (image_name)`
- `build -t (image_name:image_tag)` option: build a new image with name image_name and tag image_tag
- `prune` command. Remove all dangling images.
- `prune -a` command. Remove all unused images.
- `$(docker images -q)` = all images. Can be used with `rm`
 - e.g. `docker image rm $(docker images -q)`: removes all images

docker container

- Manage containers
- `ls` command: list. List all the running containers on the host.
 - Same as `docker ps` command
- `ls -a` command: list all. List all the containers on the host.
- `rm (container_name/id)`: remove. Remove container with name or id (container_name/id)
 - Same as `docker rm (container_name/id)`
- `exec (command) command`: execute. Run the command (command) inside the container (must be running).
 - Can be used with the option `-it` to run the command (command) interactively in the container.
- `stop (container_name/id) command`: stops a running container with name or id (container_name/id)
- `start (container_name/id) command`: starts a running container with name or id (container_name/id)
- `restart (container_name/id) command`: restarts a running container with name or id (container_name/id)
- `$(docker container ls -aq) := all containers. Can be used with rm, start, stop, restart`
 - e.g. `docker container stop $(docker container ls -aq)`: stops all containers
- `logs (container_name/id) command`: shows the logs for the container with name or id (container_name/id)
- `top (container_name/id) command`: shows the running processes int the container with name or id (container_name/id)
- `prune` command: removes all stopped containers

docker volume

- `create (volume_name) command`: creates a volume with name (volume_name)
- `ls` command: list all volumes on the host
- `rm (volume_name) command`: removes volume with name (volume_name)

docker system

- `prune` command: stops all...
 - all stopped containers
 - all networks not used by at least one container
 - all dangling images
 - all dangling build cache
- `prune -a` command: stops all...
 - all stopped containers
 - all networks not used by at least one container
 - all images without at least one container associated to them
 - all build cache